SRIVATSAN GANESH

Senior DevOps & Site Reliability Engineer (AWS, Kubernetes | AI/ML Infra)

PROFESSIONAL SUMMARY

Results-driven DevOps & Site Reliability Engineer with 4+ years of hands-on experience building, automating, and scaling high-performance cloud-native platforms. Expert in managing Kubernetes clusters for GPU/CPU workloads used in AI/ML model training and inference. Proven record in infrastructure automation, CI/CD, incident response, and HA/DR strategy implementation. Strong expertise in AWS, Terraform, observability systems, and driving platform standardization. Adept at improving uptime, optimizing deployment pipelines, and enabling rapid innovation in mission-critical environments.

PROFESSIONAL EXPERIENCE

Cloud & Platform Engineer, 02/2025 to Current
Freelance - Upwork - Remote

- Architected Kubernetes-based deployment for AI inference workloads, achieving 38% latency reduction.
- Built automated infra provisioning workflows using Terraform and Python, improving deployment speed by 45%.
- Designed and implemented centralized logging using Elasticsearch and Fluentbit, reducing troubleshooting time by 30%.
- Enabled GPU auto-scaling based on real-time ML workload demand using custom Python scripts and CloudWatch metrics.

DevOps Engineer, 04/2022 to 12/2024
Cognizant Technology Solutions India Pvt. Ltd. - Chennai, India

- Managed AWS-based Kubernetes clusters (EKS/ECS) handling GPU & CPU ML workloads, ensuring 99.99% uptime.
- Optimized GPU performance (CUDA, NCCL, driver tuning), cutting AI training cycle time by 25%.
- Created CI/CD pipelines with GitHub Actions, Jenkins & AWS CodePipeline-reduced deployment time by 60%.
- Implemented Prometheus/Grafana monitoring and CloudWatch alerting, reducing MTTR by 35%.
- Built HA/DR strategy across multi-AZ architecture using EC2, ALB, RDS failover-resulting in zero downtime migration.
- Integrated Kafka and Redis to enhance data streaming performance for AI-based applications.
- Developed self-service deployment tools using Python scripts and Terraform modules-improving DevOps efficiency by 40%.
- Participated in incident response, root cause analysis & reliability metrics tracking (SLA, SLO conformity).

CERTIFICATIONS

AWS DevOps Professional Certification, **07/2025**

DevOps Mastery Certification - KodeKloud, **05/2023**

Applied DevOps Engineering Certification - IBM, **10/2022**

CONTACT

Phone: +919677142798
Email: srivatsan.ganesh98@gmail.com
WWW: Https://srivatsanganesh.netlify.app/
LinkedIn: Linkedin.com/in/srivatsan-g/

EDUCATION

Bachelor of Engineering, Electronics And Instrumentation Engineering, 09/2020
St.Joseph's College of Engineering - Chennai, India

CORE COMPETENCIES

- **Cloud:** AWS, Azure
- **Container & Cluster Mgmt:** Kubernetes, Docker, Helm, ArgoCD, Auto Scaling
- **Distributed Tech:** Kafka, Redis, Elasticsearch, Nginx, MySQL
- **Infrastructure Automation:** Terraform, CloudFormation, Bash, Python
- **CI/CD:** GitHub Actions, Jenkins, CodePipeline, Spinnaker, GitLab CI
- **Observability & Reliability:** Prometheus, Grafana, CloudWatch, ELK, Fluentbit, SLA/SLO
- **Networking / Security:** VPC, Service Mesh (Istio), Cilium CNI, eBPF (basic), Security Groups
- **AI Systems Optimization:** GPU Computing (CUDA, NCCL), ML infra scaling, performance tuning
- **SRE Practices:** Incident Management, DR planning (HA/DR), Postmortems, Root Cause Analysis
- **MLOps:** MLflow, Kubeflow, Amazon SageMaker, DVC, Feature Store, Model Registry
- **ML Lifecycle Automation:** Data validation, model training automation, CI/CD for ML, inference pipelines
- **Cloud Platforms:** AWS (EKS, S3, EC2, IAM, RDS, Lambda), Azure
- **IaC:** Terraform, CloudFormation, Helm
- **Containers & Orchestration:** Docker, Kubernetes, EKS, ArgoCD
- **Monitoring:** Prometheus, Grafana, CloudWatch, ELK, Drift Monitoring
- **Workflow Orchestration:** Apache Airflow, Step Functions
- **Scripting:** Python, Bash
- **Specialization:** GPU computing (CUDA, NCCL)